

Network Intrusion Detection

Sofiane Alhoz

Overview

Goal :

Implement and evaluate the accuracy of different machine learning models for classification tasks using K-Fold Cross-Validation.

- Dataset
- For each model:
 - Perceptron
 - NN3
 - NN5
 - CNN2
 - CNN5

We will tackle those points: Model view → Hyperparameter Tuning → K-Fold evaluation

- Conclusion

Dataset

Network intrusion detection involves identifying network flows that carry malicious traffic.

A flow is defined as a sequence of packets sharing the same attributes, such as:

Source IP, Destination IP, Source Port,

Destination Port, Protocol

Flow features refer to the statistical characteristics extracted from these flows.

76 features per file, 5 files

Total: 2830749 samples

net_intrusion_detection > MachineLearningCVE >  Tuesday-WorkingHours.pcap_ISCX.csv

```
161976
161977
161978      ,0,9,100,15,188,29200,227,6,32,0,0,0,0,0,0,0,FTP-Patator
161979
161980
161981      ,9,105,15,188,29200,227,6,32,0,0,0,0,0,0,0,FTP-Patator
161982      0,0,9,96,15,188,29200,227,6,32,0,0,0,0,0,0,0,FTP-Patator
161983      ,0,0,0,0,9,96,15,188,29200,227,6,32,0,0,0,0,0,0,0,FTP-Patator
161984      0,0,0,9,93,15,188,29200,227,6,32,0,0,0,0,0,0,0,FTP-Patator
161985
161986      ,94,15,188,29200,227,6,32,0,0,0,0,0,0,0,0,FTP-Patator
161987
161988
161989      ,96,29200,227,3,32,0,0,0,0,0,0,0,0,0,FTP-Patator
161990      88,290,272,2,32,0,0,0,0,0,0,0,0,0,0,BENIGN
161991
161992      27.6111111,488,0,0,0,0,0,0,15,1496,18,2297,29200,257,10,32,0,0,0,0,0,0,SSH-Patator
161993
161994
161995
161996
161997      ,780,0,0,0,0,0,0,24,871,33,42384,65535,122,6,32,0,0,0,0,0,0,0,BENIGN
161998      0,0,0,0,0,19,871,30,39622,65535,122,6,32,0,0,0,0,0,0,0,0,BENIGN
161999
162000
162001
162002
162003      7,876,0,0,0,0,0,0,27,871,60,86533,65535,122,6,32,0,0,0,0,0,0,0,BENIGN
162004
162005
```



Dataset

- Emulated Data
- Imbalanced data
- 76 features (flow statistics)

	Flow Type	Number of flows
1	BENIGN	2,273,097
2	DoS Hulk	231,073
3	PortScan	158,930
4	DDoS	128,027
5	DoS GoldenEye	10,293
6	FTP-Patator	7,938
7	SSH-Patator	5,897
8	DoS slowloris	5,796
9	DoS Slowhttptest	5,499
10	Bot	1,966
11	Web Attack Brute Force	1,507
12	Web Attack XSS	652
13	Infiltration	36
14	Web Attack Sql Injection	21
15	Heartbleed	11
	Total	2,830,743
		4

Balance data:

balancing the data helps provide equal representation for all classes, preventing bias

Normalize data:

normalizing the data standardizes the dataset, improving training stability and efficiency

```
48 #We balance data as follows:
49 #1) oversample small classes so that their population/count is equal to mean_number_of_samples_per_class
50 #2) undersample large classes so that their count is equal to mean_number_of_samples_per_class
51 def balance_data(X,y,seed):
52     np.random.seed(seed)
53     unique,counts = np.unique(y,return_counts=True)
54     mean_samples_per_class = int(round(np.mean(counts)))
55     N,D = X.shape #(number of examples, number of features)
56     new_X = np.empty((0,D))
57     new_y = np.empty((0),dtype=int)
58     for i,c in enumerate(unique):
59         temp_x = X[y==c]
60         indices = np.random.choice(temp_x.shape[0],mean_samples_per_class) # gets `mean_samples_per_class` indices of class `c`
61         new_X = np.concatenate((new_X,temp_x[indices]),axis=0) # now we put new data into new_X
62         temp_y = np.ones(mean_samples_per_class,dtype=int)*c
63         new_y = np.concatenate((new_y,temp_y),axis=0)
64
65     # in order to break class order in data we need shuffling
66     indices = np.arange(new_y.shape[0])
67     np.random.shuffle(indices)
68     new_X = new_X[indices,:]
69     new_y = new_y[indices]
70     return (new_X,new_y)
71
72 # normalization: Features are normalized to have a mean close to zero and comparable variations.
73 def normalize(data):
74     data = data.astype(np.float32)
75
76     eps = 1e-15
77
78     mask = data==-1
79     data[mask]=0
80     mean_i = np.mean(data,axis=0)
81     min_i = np.min(data,axis=0) # to leave -1 (missing features) values as is and exclude in normalizing
82     max_i = np.max(data,axis=0)
83
84     r = max_i-min_i+eps
85     data = (data-mean_i)/r # zero centered
86
87     #deal with missing features -1
88     data[mask] = 0
89     return data
```

Perceptron Classifier

```
13 class Perceptron(nn.Module):
14     """
15     Perceptron Regression Model
16     """
17     def __init__(self, input_dim, num_classes, device):
18         | super(Perceptron, self).__init__()
19         | self.classifier = nn.Linear(input_dim, num_classes).to(device)
20
21     def forward(self, x):
22         | output= self.classifier(x)
23         | return output
```

Then I tried to identify the best combinations of hyperparameters from a predefined set of values, specifically for **learning rate** and **regularization strength**. These hyperparameters play a crucial role in determining the **convergence speed** and the model's ability to **generalize without overfitting**. By testing different configurations, we can refine the model for better performance in later runs.

Perceptron Classifier

Configurations:

1. Batch_size = 1024
2. learning_rates = [1e-4, 1e-2, 1e-1]
3. regularizations = [1e-6, 1e-4, 1e-3]
4. Epoch = 25


Training Accuracy on above hyperparameters:

```
[26] |  
...  
(0.0001, 1e-06) -> 60.83  
(0.0001, 0.0001) -> 60.38  
(0.0001, 0.01) -> 47.48  
(0.01, 1e-06) -> 80.50  
(0.01, 0.0001) -> 66.94  
(0.01, 0.01) -> 54.83  
(1.0, 1e-06) -> 78.96  
(1.0, 0.0001) -> 67.60  
(1.0, 0.01) -> 50.63
```

Results:

Best Learning rate is: [1e-1]

Best Regularization rate is : [1e-6]



KFold Evaluation

Split train 80% / test 20% might be problematic because attacks happens at different and precise moments

=> We need to use KFold Evaluation:

The model is trained **K times**, each time using **K-1 folds for training** and the **remaining fold for testing**.

5Fold Evaluation:

Itération 1: [Train | Train | Train | Train | Test]

Itération 2: [Train | Train | Train | Test | Train]

Itération 3: [Train | Train | Test | Train | Train]

Itération 4: [Train | Test | Train | Train | Train]

Itération 5: [Test | Train | Train | Train | Train]

This ensures that every data point is used for both training and testing, leading to **better generalization** and reducing the risk of **overfitting**.

```
Fold #0
INFO:models:Classifier ini
INFO:models:Starting train
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [1/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [2/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
DEBUG:models:Epoch [3/20],
...
DEBUG:models:Epoch [9/20],
DEBUG:models:Epoch [9/20],
DEBUG:models:Epoch [10/20]
WARNING:models:No improvem
```

Output is truncated. View as a [scroll](#)
Loaded MachineLearningCVE/
balanced test acc: 73.825

Perceptron Classifier

Fold #0

balanced test acc: 73.82532672957981

Fold #1

balanced test acc: 79.07317125426042

Fold #2

balanced test acc: 77.87223704729863

Fold #3

balanced test acc: 81.22136833071558

Fold #4

balanced test acc: 69.37087913280486

76.27

NN3 architecture


```
107 class Net3(nn.Module):
108     """
109     A neural network model consisting of multiple layers, used for classification tasks.
110
111     Args:
112     - input_dim (int): The dimensionality of the input data.
113     - num_classes (int): The number of classes in the classification problem.
114     - device (str): The device to be used for running the computations.
115
116     Attributes:
117     - input_dim (int): The dimensionality of the input data.
118     - num_classes (int): The number of classes in the classification problem.
119     - device (str): The device to be used for running the computations.
120
121     Methods:
122     - forward(x): Defines the forward pass of the model, taking in an input tensor x and returning the output tensor.
123
124     """
125     def __init__(self, input_dim, num_classes, device):
126         super(Net3, self).__init__()
127         # kernel
128         print('building NN3')
129         self.input_dim = input_dim
130         self.num_classes = num_classes
131
132         layers = []
133         layers.append(nn.Dropout(p=0.1))
134         layers.append(nn.Linear(self.input_dim, 128))
135         layers.append(nn.BatchNorm1d(num_features=128))
136         layers.append(nn.Dropout(p=0.3))
137         layers.append(nn.Linear(128, 128))
138         layers.append(nn.BatchNorm1d(num_features=128))
139         layers.append(nn.Linear(128, self.num_classes))
140         self.classifier = nn.Sequential(*layers).to(device)
141
142     def forward(self, x):
143         x = self.classifier(x)
144         return x
```

NN3: Hyperparameter Tuning 1st run

Name	Smoothed	Value	Step	Time	Relative
optim_Adam_lr_0.01_reg_1e-07_bs_1024	0.1355	0.1199	20.39M	Wed Nov 20, 01:13:41	6m 46s
optim_Adam_lr_0.003_reg_1e-07_bs_1024	0.1368	0.1307	20.39M	Wed Nov 20, 00:52:19	6m 46s
optim_Adam_lr_0.001_reg_1e-07_bs_1024	0.1426	0.1375	20.39M	Wed Nov 20, 00:30:59	6m 47s
optim_Adam_lr_0.003_reg_1e-05_bs_1024	0.1514	0.1471	20.39M	Wed Nov 20, 00:59:26	6m 46s
optim_Adam_lr_0.001_reg_1e-05_bs_1024	0.1553	0.1511	20.39M	Wed Nov 20, 00:38:05	6m 44s
optim_Adam_lr_0.01_reg_1e-05_bs_1024	0.1561	0.1519	20.39M	Wed Nov 20, 01:20:47	6m 44s
optim_Adam_lr_0.0003_reg_1e-07_bs_1024	0.1773	0.1735	20.39M	Wed Nov 20, 00:09:35	6m 45s
optim_Adam_lr_0.0003_reg_1e-05_bs_1024	0.1939	0.2188	20.39M	Wed Nov 20, 00:16:42	6m 45s
optim_Adam_lr_0.0001_reg_1e-07_bs_1024	0.2091	0.1748	20.39M	Tue Nov 19, 23:48:17	6m 45s
optim_Adam_lr_0.0001_reg_1e-05_bs_1024	0.224	0.2257	20.39M	Tue Nov 19, 23:55:21	6m 43s
optim_Adam_lr_0.001_reg_0.001_bs_1024	0.3177	0.2829	20.39M	Wed Nov 20, 00:45:12	6m 45s
optim_Adam_lr_0.0003_reg_0.001_bs_1024	0.3214	0.2659	20.39M	Wed Nov 20, 00:23:50	6m 46s
optim_Adam_lr_0.003_reg_0.001_bs_1024	0.3277	0.2992	20.39M	Wed Nov 20, 01:06:33	6m 45s
optim_Adam_lr_0.01_reg_0.001_bs_1024	0.3316	0.3406	20.39M	Wed Nov 20, 01:27:52	6m 44s
optim_Adam_lr_0.0001_reg_0.001_bs_1024	0.3503	0.3396	20.39M	Wed Nov 20, 00:02:29	6m 46s


Settings:

1. learning_rates =
[1e-4, 3e-4, 1e-3, 3e-3, 1e-2]
2. regularizations = [1e-7, 1e-5, 1e-3]



Results:

1. Need more epochs: Because, as we can see for example in the loss curves, the model has not yet completed its learning process.
2. plot: Best val accuracy: 88.74
3. plot : Best parameters: (1e-3, 1e-5)



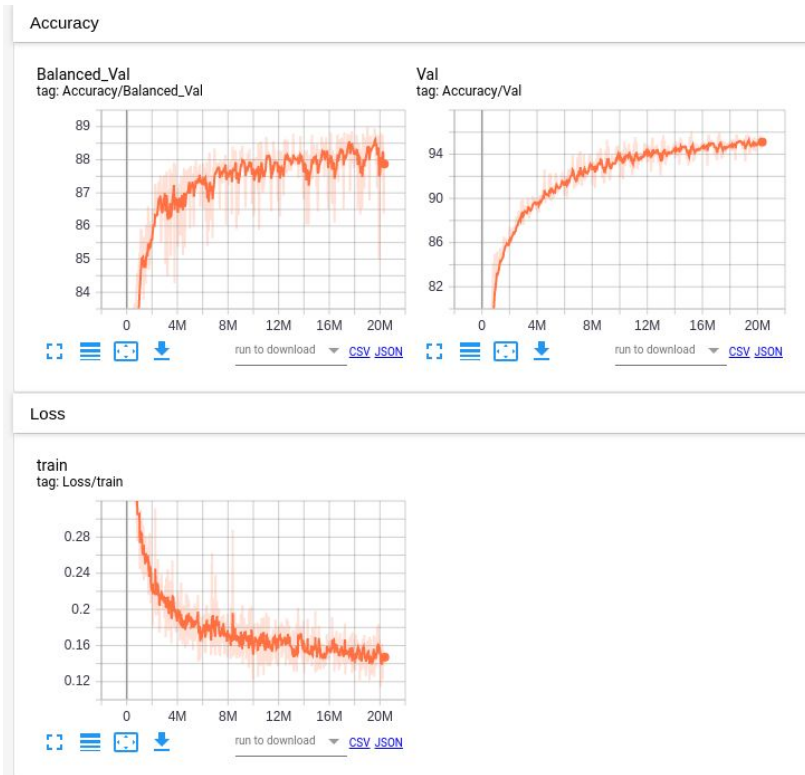
NN3: 2nd run with batchnorm

Settings:

1. learning_rates = [1e-3]
2. regularizations = [1e-5]
3. Num_epochs = 10
4. **Use of batchnorm**

Results:

1. Best val accuracy: 88.98, small improvement
2. Graph feedback: It is still learning. Need more epochs



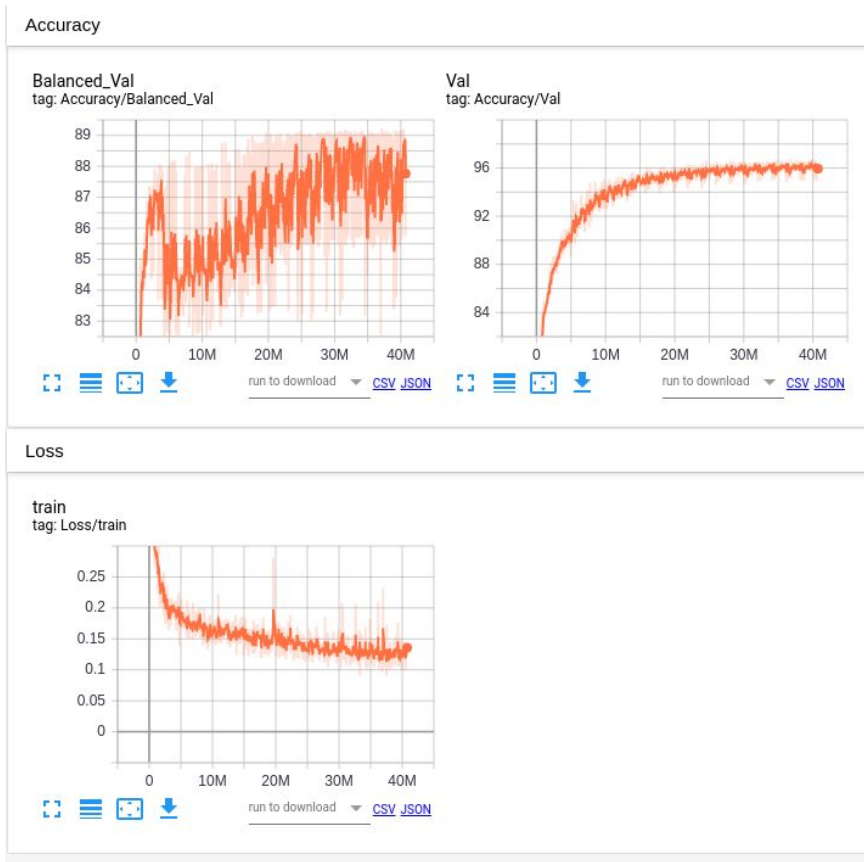
NN3: 3rd run with batchnorm + dropout

Settings:

1. learning_rates = [1e-3]
2. regularizations = [1e-5]
3. Use of batchnorm
4. **Use of dropout (0.1,0.3)**
5. **num_epochs=20**

Results:

1. Best val accuracy: 89.22, small improvement



NN3: 5-Fold evaluation

Settings:

1. `learning_rates = [1e-3]`
2. `regularizations = [1e-5]`
3. Use batchnorm
4. Use dropout (0.1,0.3)
5. `num_epochs=20`

Results of the 5 fold::

1. Balanced Test acc: 85.61

Fold #4

INFO:models:Classifier initialized with method nn3, input_dim 76, num_classes 15, num_epochs 20

INFO:models:Starting training process...

building NN3

DEBUG:models:Epoch [1/20], Step [50/398], Loss: 1.0672

DEBUG:models:Epoch [1/20], Step [100/398], Loss: 0.6369

DEBUG:models:Epoch [1/20], Step [150/398], Loss: 0.5170

DEBUG:models:Epoch [1/20], Step [200/398], Loss: 0.4320

DEBUG:models:Epoch [1/20], Step [250/398], Loss: 0.4075

DEBUG:models:Epoch [1/20], Step [300/398], Loss: 0.3748

DEBUG:models:Epoch [1/20], Step [350/398], Loss: 0.3616

DEBUG:models:Epoch [2/20], Step [1/398], Loss: 0.3724

DEBUG:models:Epoch [2/20], Step [51/398], Loss: 0.3449

DEBUG:models:Epoch [2/20], Step [101/398], Loss: 0.3204

DEBUG:models:Epoch [2/20], Step [151/398], Loss: 0.3445

DEBUG:models:Epoch [2/20], Step [201/398], Loss: 0.3468

DEBUG:models:Epoch [2/20], Step [251/398], Loss: 0.3178

DEBUG:models:Epoch [2/20], Step [301/398], Loss: 0.3099

DEBUG:models:Epoch [2/20], Step [351/398], Loss: 0.2928

DEBUG:models:Epoch [3/20], Step [2/398], Loss: 0.3019

DEBUG:models:Epoch [3/20], Step [52/398], Loss: 0.2696

DEBUG:models:Epoch [3/20], Step [102/398], Loss: 0.2973

DEBUG:models:Epoch [3/20], Step [152/398], Loss: 0.2940

DEBUG:models:Epoch [3/20], Step [202/398], Loss: 0.2739

DEBUG:models:Epoch [3/20], Step [252/398], Loss: 0.2703

DEBUG:models:Epoch [3/20], Step [302/398], Loss: 0.2526

DEBUG:models:Epoch [3/20], Step [352/398], Loss: 0.2712

DEBUG:models:Epoch [4/20], Step [3/398], Loss: 0.2667

DEBUG:models:Epoch [4/20], Step [53/398], Loss: 0.2632

...

DEBUG:models:Epoch [7/20], Step [306/398], Loss: 0.2246

DEBUG:models:Epoch [7/20], Step [356/398], Loss: 0.2239

DEBUG:models:Epoch [8/20], Step [7/398], Loss: 0.2439

WARNING:models:No improvement in accuracy for 10 iterations.

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

Loaded MachineLearningCVE/runs/nn3/Kfold-4th_run/optim_Adam_lr_0.001_reg_1e-05_bs_5120

balanced test acc: 86.24194198056313

85.73

5 layer NN architecture

```
147 class Net5(nn.Module):
148     """
149     A neural network model with 4 fully connected layers, using batch normalization and dropout for regularization.
150     Args:
151     - input_dim (int): the number of input features
152     - num_classes (int): the number of output classes
153     - device (torch.device): the device on which to run the model
154     Attributes:
155     - input_dim (int): the number of input features
156     - num_classes (int): the number of output classes
157     - model (nn.Sequential): the neural network architecture consisting of 4 fully connected layers
158     """
159     def __init__(self, input_dim, num_classes, device):
160         super(Net5, self).__init__()
161         # kernel
162         self.input_dim = input_dim
163         self.num_classes = num_classes
164         layers = []
165         layers.append(nn.Linear(input_dim, 128))
166
167         layers.append(nn.BatchNorm1d(128))
168         layers.append(nn.ReLU(True))
169         layers.append(nn.Linear(128, 256))
170
171         layers.append(nn.BatchNorm1d(256))
172         layers.append(nn.Dropout(p=0.3))
173         layers.append(nn.ReLU(True))
174         layers.append(nn.Linear(256, 256))
175
176         layers.append(nn.BatchNorm1d(256))
177         layers.append(nn.Dropout(p=0.4))
178         layers.append(nn.ReLU(True))
179         layers.append(nn.Linear(256, 128))
180
181         layers.append(nn.BatchNorm1d(128))
182         layers.append(nn.Dropout(p=0.5))
183         layers.append(nn.ReLU(True))
184         layers.append(nn.Linear(128, num_classes))
185
186         self.model = nn.Sequential(*layers).to(device)
187     def forward(self, x):
188         return self.model(x)
```


NN5: Hyperparameter Tuning: 1st run: Dropout higher layers

Configurations:


1. learning_rates =
[3e-2, 1e-3, 3e-4, 1e-4]
2. regularizations = [1e-6, 1e-5, 1e-4]
3. Batchnorm
4. Dropout prob: (0, 0, .3, .4, .5)
5. Epoch = 20

Results:

1. Val acc: 90.67

Feedback:

1. Bigger reg
2. Bigger lr
3. Larger epoch



NN5: Run 2: Dropout higher layers

Configurations:

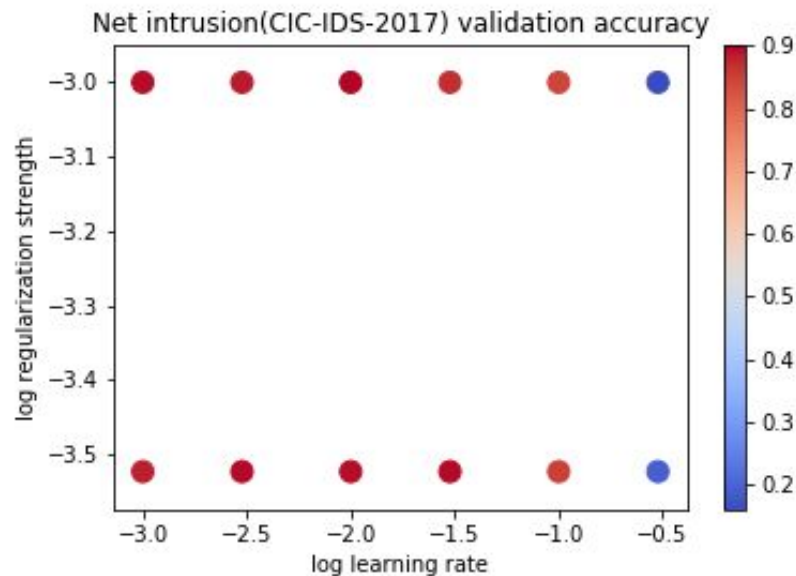
1. learning_rates = [1e-3, 3e-3, 1e-2, 3e-2, 1e-1, 3e-1]
2. regularizations = [3e-4, 1e-3]
3. Batchnorm
4. Dropout prob: (0, 0, .3, .4, .5)
5. Epoch = 60

Results:

1. Val acc: 90.27

Feedback:

1. Search does not have to be dense
2. Bigger reg range
3. smaller lr
4. Conflicting to previous run, it could be because of small epoch number previously



NN5: Run 3: Dropout higher layers

Configurations:


1. learning_rates = [1e-5, 1e-4, 3e-4]
2. regularizations = [1e-6, 1e-5, 1e-4, 1e-3]
3. Batchnorm
4. Dropout prob: (0, 0, .3, .4, .5)
5. Epoch = 60

Results:

1. Val acc: 91.14,

Feedback:

1. Best param lr, reg = (1e-4, 1e-5)



5-fold Evaluation: NN5

Settings

1. Best param $lr, reg = (1e-4, 1e-5)$
2. Epochs = 60

Results:

1. 5-Fold Balanced Test acc: **85.63**

CNN2 architecture

```
26 class CNN2(nn.Module):
27     """
28     2-layer Convolutional Neural Network Model
29     """
30     def __init__(self, input_dim, num_classes, device):
31         super(CNN2, self).__init__()
32         # kernel
33         self.input_dim = input_dim
34         self.num_classes = num_classes
35
36         conv_layers = []
37         conv_layers.append(nn.Conv1d(in_channels=1, out_channels=64, kernel_size=3, padding=1)) # ;input_dim, 64
38         conv_layers.append(nn.BatchNorm1d(64))
39         conv_layers.append(nn.ReLU(True))
40
41         conv_layers.append(nn.Conv1d(in_channels=64, out_channels=128, kernel_size=3, padding=1)) #(input_dim, 128)
42         conv_layers.append(nn.BatchNorm1d(128))
43         conv_layers.append(nn.ReLU(True))
44
45         self.conv = nn.Sequential(*conv_layers).to(device)
46
47         fc_layers = []
48         fc_layers.append(nn.Linear(input_dim*128, num_classes))
49         self.classifier = nn.Sequential(*fc_layers).to(device)
50
51     def forward(self, x):
52         batch_size, D = x.shape
53         x = x.view(batch_size, 1, D)
54
55         x = self.conv(x)
56         x = torch.flatten(x, 1)
57         x = self.classifier(x)
58         return x
```

CNN2: run 1

Config::

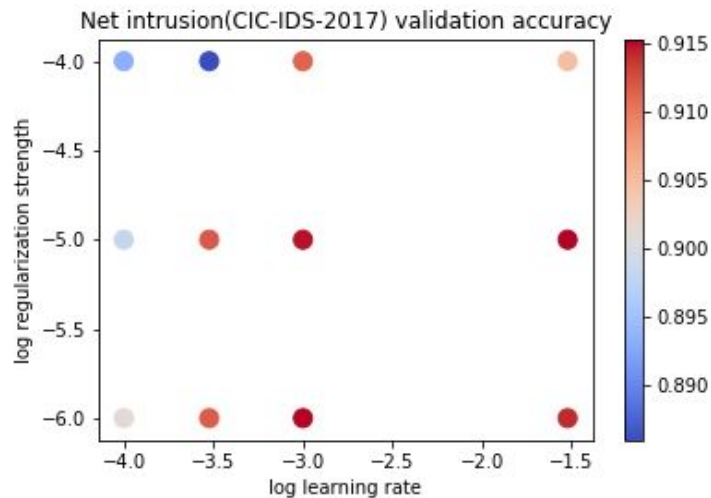
1. num_epochs = 60
2. learning_rates = [3e-2, 1e-3, 3e-4, 1e-4]
3. regularizations = [1e-6, 1e-5, 1e-4]

acc:

1. Val:91.53

Feedback:

1. Lr 1e-3 is best, maybe need smaller reg



CNN2: run 2

Configs:


1. num_epochs = 60
2. learning_rates = [1e-3]
3. regularizations = [1e-8, 1e-7, 1e-6]

acc:

1. Val:91.76

Feedback:

1. Best reg is 1e-6



5-fold Evaluation: CNN2

Settings:

1. `num_epochs = 60`
2. `learning_rates = [1e-3]`
3. `regularizations = [1e-6]`

Results:

1. **5-Fold acc: 88.14%**

CNN5 architecture

```
60 class CNN5(nn.Module):
61     """
62     5-layer Convolutional Neural Network Model
63     """
64     def __init__(self, input_dim, num_classes, device):
65         super(CNN5, self).__init__()
66         # kernel
67         self.input_dim = input_dim
68         self.num_classes = num_classes
69
70         conv_layers = []
71         conv_layers.append(nn.Conv1d(in_channels=1, out_channels=64, kernel_size=3, padding=1)) # ;input_dim,64
72         conv_layers.append(nn.BatchNorm1d(64))
73         conv_layers.append(nn.ReLU(True))
74
75         conv_layers.append(nn.Conv1d(in_channels=64, out_channels=128, kernel_size=3, padding=1)) #(input_dim,128)
76         conv_layers.append(nn.BatchNorm1d(128))
77         conv_layers.append(nn.ReLU(True))
78
79         conv_layers.append(nn.Conv1d(in_channels=128, out_channels=256, kernel_size=3, padding=1)) #(input_dim,128)
80         conv_layers.append(nn.BatchNorm1d(256))
81         conv_layers.append(nn.ReLU(True))
82
83         conv_layers.append(nn.Conv1d(in_channels=256, out_channels=256, kernel_size=3, padding=1)) #(input_dim,128)
84         conv_layers.append(nn.BatchNorm1d(256))
85         conv_layers.append(nn.ReLU(True))
86
87         conv_layers.append(nn.Conv1d(in_channels=256, out_channels=128, kernel_size=3, padding=1)) #(input_dim,128)
88         conv_layers.append(nn.BatchNorm1d(128))
89         conv_layers.append(nn.ReLU(True))
90
91         self.conv = nn.Sequential(*conv_layers).to(device)
92         fc_layers = []
93         fc_layers.append(nn.Linear(input_dim*128, num_classes))
94         self.classifier = nn.Sequential(*fc_layers).to(device)
95
96     def forward(self, x):
97         batch_size, D = x.shape
98         x = x.view(batch_size, 1, D)
99         x = self.conv(x)
100         x = torch.flatten(x, 1)
101         x = self.classifier(x)
102         return x
```

5-fold Evaluation: CNN5

Settings:

1. num_epochs = 60
2. learning_rates = [1e-3]
3. regularizations = [1e-6]

Results:

1. **5-Fold** acc: 88.18%

Conclusion

Classifier	5-fold Accuracy
Perceptron	76.27
NN3	85.73
NN5	85.61
CNN2	88.14
CNN5	88.18

CNN5 achieved the best results

The Perceptron results confirm that **linear classification is insufficient** for complex intrusion detection tasks.

Also, the results shows that depth increase doesn't make a big difference, contrary to the change from NN to CNN, which makes a big difference.